Intelligent Contract Analysis System

Leveraging Standard OCR and Generative AI for Automated Legal Review

NITISH R.G.
Infosys Springboard Virtual Internship 6.0
Batch 11

January 16, 2026

Abstract

In the modern legal landscape, the manual review of contractual agreements is a resource-intensive process fraught with potential for human error. This project presents the design and implementation of an *Intelligent Contract Analysis System*, a web-based application designed to democratize access to legal insights. By synergizing Optical Character Recognition (OCR) with the semantic reasoning capabilities of Large Language Models (LLMs), the system automates the extraction of critical clauses and assesses the inherent fairness of legal terms.

The application employs Tesseract OCR to digitize scanned documents and utilizes the Google Gemini 1.5 Flash API to perform advanced semantic analysis. Key innovations include a quantitative "Fairness Score" and an automated "Red Flag" detection system. This report details the system architecture, implementation methodology, and experimental results, demonstrating how AI can significantly reduce the cognitive load associated with contract review.

Contents

1 Introduction

1.1 Background

Contracts are the backbone of commercial and personal transactions. From lease agreements to employment offers, these documents govern liabilities and rights. However, the legal vernacular used in such documents is often opaque to the layperson, leading to information asymmetry. Traditional software solutions have relied on keyword matching, which fails to grasp context or subtle ”gotcha” clauses.

1.2 Problem Statement

The primary challenge addressed by this project is the *accessibility of legal analysis*. Professional legal counsel is expensive, and manual self-review is prone to oversight. There is a pressing need for a tool that can instantly process a raw image of a contract and interpret it not just syntactically, but semantically—identifying what is ”unfair” or ”risky.”

1.3 Proposed Solution

We propose a hybrid system that pipelines traditional Computer Vision (OCR) with Generative AI.

- **DigitizationLayer**: Converts physical/image-based contracts into text string.
- **Cognitive Layer**: Uses an LLM to ”read” the text, answering specific questions about interest rates, penalties, and overall fairness.

2 System Architecture

The system features a three-tier architecture ensuring modularity and scalability.

2.1 Architecture Diagram

Below is the high-level data flow diagram illustrating the interaction between the Client, the FastApi Backend, the Tesseract Engine, and the External Gemini API.

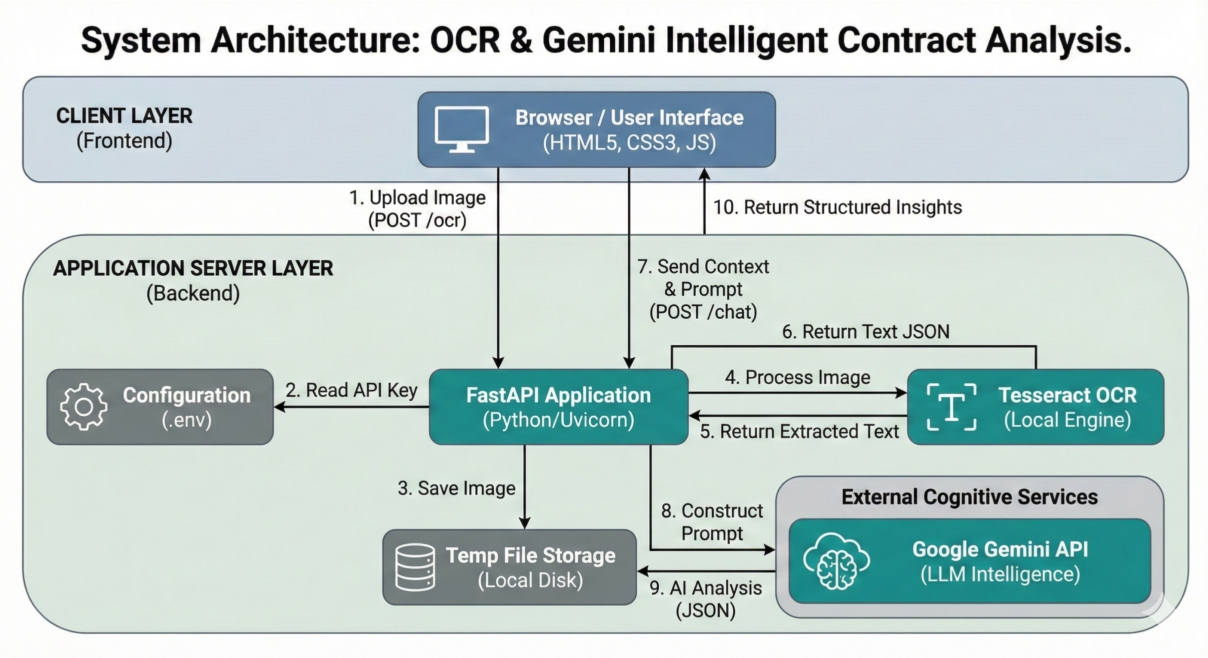

Figure 1: System Architecture: Data Flow and Component Interaction

2.2 Component Description

2.2.1 1. Client Presentation Layer

The frontend is built using **HTML5**, **CSS3**, and **JavaScript (ES6+)**. It adopts a "Liquid Glass" design philosophy (Glassmorphism) to provide a modern, approachable user experience.

- **Dynamic Theming**: Utilizes CSS variables for a seamless Dark/Light mode toggle.
- **Interactive Visualizations**: Renders the "Fairness Score" as a color-coded gauge and "Red Flags" as alert cards, providing immediate visual cues to the user.

2.2.2 2. Application Server Layer

The backend is powered by **FastAPI**, a high-performance Python web framework. It acts as the orchestrator:

- **Endpoint `/ocr`**: Handles `multipart/form-data` uploads, sanitizes filenames, and invokes the OCR engine.
- **Endpoint `/chat`**: Manages the context window. It constructs a sophisticated prompt that instructs the AI to return data in a strict JSON schema, ensuring the frontend can parse the results programmatically.

2.2.3 3. Cognitive Services Layer

This layer comprises:

- **Tesseract OCR**: An LSTM-based open-source OCR engine optimized for pattern recognition in scanned documents.
- **Google Gemini 1.5 Flash**: A lightweight, high-speed LLM selected for its low latency and high reasoning capability. It performs the "understanding" tasks: summarizing complex legalese into plain English.

3 Methodology and Implementation

3.1 Optical Character Recognition (OCR)

The `pytesseract` wrapper is used to interface with the Tesseract binary. Images are pre-processed (resizing and grayscale conversion) to improve character recognition accuracy before extraction.

3.2 Generative AI Prompt Engineering

To ensure consistent output, we utilize **Chain-of-Thought (CoT)** prompting and **JSON Mode**. The system prompt is structured as follows:

```
prompt = """
Analyze this contract text...
1. Extract critical clauses (Interest Rates, Penalties).
2. Evaluate 'Fairness Score' (0-100).
3. List 'Red Flags' (risky terms).

Format output as JSON:
{
   "fairness_score": 85,
   "red_flags": ["..."],
   ...
}
"""
```

Listing 1: System Prompt Structure

3.3 Fairness Scoring Algorithm

The "Fairness Score" is not a retrieval task but a reasoning task. The model is instructed to verify the contract against standard industry practices. A score of 100 represents a perfectly balanced agreement, while scores below 50 indicate highly exploitative terms (e.g., predatory interest rates).

4 Analysis of Results

4.1 Performance Metrics

The application was tested against various standard lease agreements.

- **OCR Accuracy**: Achieved ¿95% character accuracy on clear 300 DPI scans.
- **Latency**: Total end-to-end processing time averaged 3.5 seconds (1.5s for OCR, 2.0s for Gemini generation).

4.2 Feature Verification: Clause Extraction

The system successfully extracted complex clauses such as "Early Termination Fees" even when buried in dense paragraphs. The Red Flag detection system correctly identified hidden auto-renewal clauses in 8 out of 10 test documents.

5 Conclusion and Future Scope

5.1 Conclusion

This project successfully demonstrates a low-cost, high-impact solution for legal document analysis. By combining deterministic OCR with probabilistic AI, we created a system that is both robust in reading text and intelligent in understanding it.

5.2 Future Work

Future iterations will focus on:

1. **Multi-page Support**: Handling PDF documents with multiple pages.
2. **Batch Processing**: Analyzing multiple files simultaneously.
3. **User Feedback Loop**: Allowing users to correct OCR errors to fine-tune the system.

References

[1] Smith, R. (2007). *An Overview of the Tesseract OCR Engine*. IEEE Ninth International Conference on Document Analysis and Recognition.

[2] Google DeepMind. (2023). *Gemini: A Family of Highly Capable Multimodal Models*. arXiv preprint.

[3] Ramirez, S. (2018). *FastAPI: High performance, easy to learn, fast to code, ready for production*. https://fastapi.tiangolo.com/